

1. Download the dataset: Dataset

2. Load the dataset.

```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
df=pd.read_csv(r'Churn_Modelling.csv')
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

3. Perform the Below Visualizations.


● Univariate Analysis ● Bi - Variate Analysis ● Multivariate Analysis

```
sns.distplot(df['Age'])
```

```
C:\Users\AdhoLOKHAM\anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
```

```
code to use either `displot` (a figure-level function with similar
flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)
```


```
<AxesSubplot:xlabel='Age', ylabel='Density'>
```



```
sns.lineplot(df['Age'],df['Exited'])
```


```
C:\Users\AdhoLOKHAM\anaconda3\lib\site-packages\seaborn\
_decorators.py:36: FutureWarning: Pass the following variables as
keyword args: x, y. From version 0.12, the only valid positional
argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.
warnings.warn(
```

```
<AxesSubplot:xlabel='Age', ylabel='Exited'>
```



```
plt.pie(df.NumOfProducts.value_counts(),colors=['red','green'],autopct
='%.1f%%')
plt.title('NumOfProducts')
Text(0.5, 1.0, 'NumOfProducts')
```


NumOfProducts



```
sns.barplot((df.Geography.value_counts()).index,df.Geography.value_cou
nts())
```


```
C:\Users\AdhoLOKHAM\anaconda3\lib\site-packages\seaborn\
_decorators.py:36: FutureWarning: Pass the following variables as
keyword args: x, y. From version 0.12, the only valid positional
argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.
  warnings.warn(
```

```
<AxesSubplot:ylabel='Geography'>
```




```
sns.scatterplot(x=df.CreditScore,y=df.Balance)
```

```
<AxesSubplot:xlabel='CreditScore', ylabel='Balance'>
```



```
plt.pie(df.Geography.value_counts(),colors=['red','green','blue'],labels=['France','Germany','Spain'],autopct='%1f%%')
plt.title('Geography')
```

```
Text(0.5, 1.0, 'Geography')
```



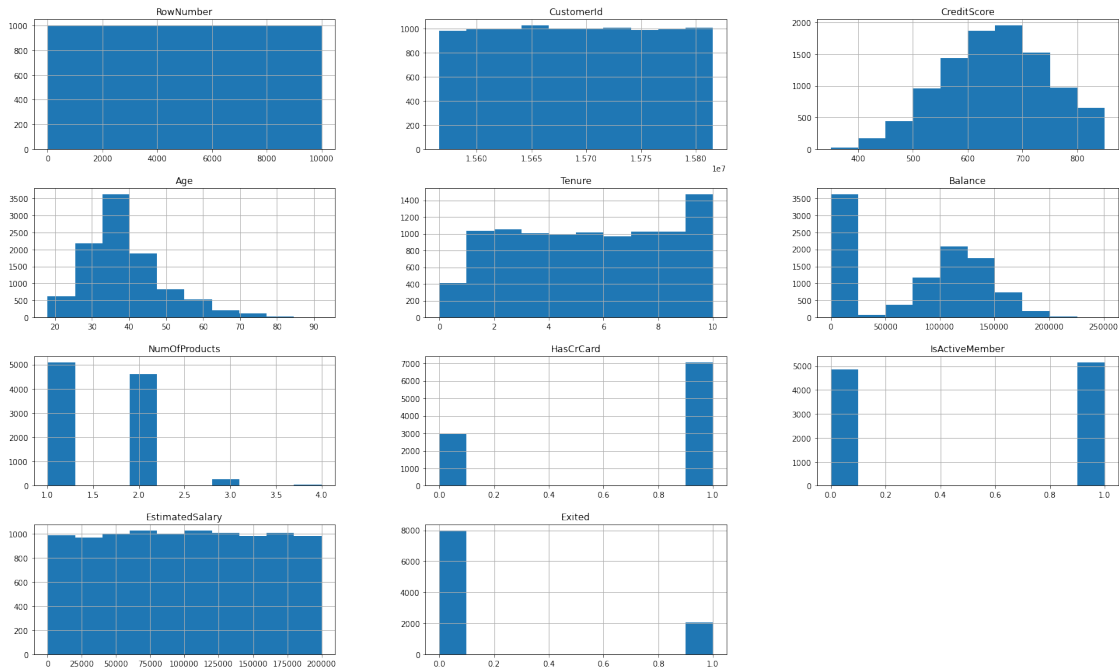
```
df.hist(figsize=(25,15))
```

```
array([[<AxesSubplot:title={'center':'RowNumber'}>,
        <AxesSubplot:title={'center':'CustomerId'}>],
```

```

<AxesSubplot:title={'center':'CreditScore'}>],
[<AxesSubplot:title={'center':'Age'}>,
<AxesSubplot:title={'center':'Tenure'}>,
<AxesSubplot:title={'center':'Balance'}>],
[<AxesSubplot:title={'center':'NumOfProducts'}>,
<AxesSubplot:title={'center':'HasCrCard'}>,
<AxesSubplot:title={'center':'IsActiveMember'}>],
[<AxesSubplot:title={'center':'EstimatedSalary'}>,
<AxesSubplot:title={'center':'Exited'}>, <AxesSubplot:>]],
dtype=object)

```




```
sns.boxplot(df.Age)
```

C:\Users\AdhoLOKHAM\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
<AxesSubplot:xlabel='Age'>
```



4. Perform descriptive statistics on the dataset.

```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age
Tenure \				
count	10000.00000	1.000000e+04	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800
std	2886.89568	7.193619e+04	96.653299	10.487806
min	1.00000	1.556570e+07	350.000000	18.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000
max	10000.00000	1.581569e+07	850.000000	92.000000
count	10000.000000	10000.000000	10000.00000	10000.000000
mean	76485.889288	1.530200	0.70550	0.515100
std	62397.405202	0.581654	0.45584	0.499797
min	0.000000	1.000000	0.00000	0.000000

25%	0.000000	1.000000	0.000000	0.000000
50%	97198.540000	1.000000	1.000000	1.000000
75%	127644.240000	2.000000	1.000000	1.000000
max	250898.090000	4.000000	1.000000	1.000000

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

```
df.mean()
```

```
C:\Users\ADH0L0~1\AppData\Local\Temp\ipykernel_2088\3698961737.py:1:
FutureWarning: Dropping of nuisance columns in DataFrame reductions
(with 'numeric_only=None') is deprecated; in a future version this
will raise TypeError. Select only valid columns before calling the
reduction.
```

```
df.mean()
```

RowNumber	5.000500e+03
CustomerId	1.569094e+07
CreditScore	6.505288e+02
Age	3.892180e+01
Tenure	5.012800e+00
Balance	7.648589e+04
NumOfProducts	1.530200e+00
HasCrCard	7.055000e-01
IsActiveMember	5.151000e-01
EstimatedSalary	1.000902e+05
Exited	2.037000e-01

dtype: float64

```
df.median()
```

```
C:\Users\ADH0L0~1\AppData\Local\Temp\ipykernel_2088\530051474.py:1:
FutureWarning: Dropping of nuisance columns in DataFrame reductions
(with 'numeric_only=None') is deprecated; in a future version this
will raise TypeError. Select only valid columns before calling the
reduction.
```

```
df.median()
```

RowNumber	5.000500e+03
CustomerId	1.569074e+07
CreditScore	6.520000e+02
Age	3.700000e+01
Tenure	5.000000e+00


```

Balance          9.719854e+04
NumOfProducts    1.000000e+00
HasCrCard         1.000000e+00
IsActiveMember    1.000000e+00
EstimatedSalary   1.001939e+05
Exited            0.000000e+00
dtype: float64

```

```
df.mode()
```

```

      RowNumber  CustomerId  Surname  CreditScore  Geography  Gender
Age \
0           1      15565701    Smith           850.0      France    Male
37.0
1           2      15565706      NaN           NaN         NaN      NaN
NaN
2           3      15565714      NaN           NaN         NaN      NaN
NaN
3           4      15565779      NaN           NaN         NaN      NaN
NaN
4           5      15565796      NaN           NaN         NaN      NaN
NaN
...         ...         ...         ...         ...         ...         ..
.
9995        9996      15815628      NaN           NaN         NaN      NaN
NaN
9996        9997      15815645      NaN           NaN         NaN      NaN
NaN
9997        9998      15815656      NaN           NaN         NaN      NaN
NaN
9998        9999      15815660      NaN           NaN         NaN      NaN
NaN
9999       10000      15815690      NaN           NaN         NaN      NaN
NaN

```

```

      Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  \
0          2.0       0.0             1.0         1.0             1.0
1          NaN       NaN             NaN         NaN             NaN
2          NaN       NaN             NaN         NaN             NaN
3          NaN       NaN             NaN         NaN             NaN
4          NaN       NaN             NaN         NaN             NaN
...         ...         ...         ...         ...             ...
9995        NaN       NaN             NaN         NaN             NaN
9996        NaN       NaN             NaN         NaN             NaN
9997        NaN       NaN             NaN         NaN             NaN
9998        NaN       NaN             NaN         NaN             NaN
9999        NaN       NaN             NaN         NaN             NaN

```

```

      EstimatedSalary  Exited
0          24924.92      0.0

```

```

1      NaN      NaN
2      NaN      NaN
3      NaN      NaN
4      NaN      NaN
...
9995   NaN      NaN
9996   NaN      NaN
9997   NaN      NaN
9998   NaN      NaN
9999   NaN      NaN

```

[10000 rows x 14 columns]

```
df.std()
```

```
C:\Users\ADH0L0~1\AppData\Local\Temp\ipykernel_2088\3390915376.py:1:
FutureWarning: Dropping of nuisance columns in DataFrame reductions
(with 'numeric_only=None') is deprecated; in a future version this
will raise TypeError.  Select only valid columns before calling the
reduction.
```

```
df.std()
```

```

RowNumber      2886.895680
CustomerId      71936.186123
CreditScore      96.653299
Age             10.487806
Tenure          2.892174
Balance        62397.405202
NumOfProducts      0.581654
HasCrCard        0.455840
IsActiveMember    0.499797
EstimatedSalary  57510.492818
Exited           0.402769
dtype: float64

```

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

```
Tenure      Balance  NumOfProducts  HasCrCard  IsActiveMember \
```

0	2	0.00	1	1	1
1	1	83807.86	1	0	1
2	8	159660.80	3	1	0
3	1	0.00	2	0	0
4	2	125510.82	1	1	1

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

5. Handle the Missing values.

```
df.isnull().sum()
```

```

RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64

```

```
df['Age'].fillna(df['Age'].mean(),inplace=True)
```

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43


	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

6. Find the outliers and replace the outliers


```
sns.displot(df['Age'])
```

```
<seaborn.axisgrid.FacetGrid at 0x1ced3490fd0>
```




```
sns.boxplot(x='Gender',y='Age',data=df)
```

```
<AxesSubplot:xlabel='Gender', ylabel='Age'>
```



```
sns.boxplot(y='Age', data=df)
```

```
<AxesSubplot:ylabel='Age'>
```




```
df['Age'].mean()
```

```
38.9218
```

```
df1=df[df['Age']<55]
```

```
sns.boxplot(y='Age',data=df1)
```

```
<AxesSubplot:ylabel='Age'>
```




```
df1['Age'].mean()
```

```
36.62250493529283
```

```
df1.boxplot()
```

```
<AxesSubplot:>
```



7. Check for Categorical columns and perform encoding.

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
df['Geography']= le.fit_transform(df['Geography'])
df.Gender= le.fit_transform(df.Gender)
df['Geography'].unique()
```

```
array([0, 2, 1], dtype=int64)
```

```
df['Gender'].unique()
```

```
array([0, 1], dtype=int64)
```

```
df=pd.get_dummies(df,columns=['Geography'])
```

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Gender	Age	
Tenure \							
0	1	15634602	Hargrave	619	Female	42	2
1	2	15647311	Hill	608	Female	41	1
2	3	15619304	Onio	502	Female	42	8
3	4	15701354	Boni	699	Female	39	1
4	5	15737888	Mitchell	850	Female	43	2

	Balance	NumOfProducts	HasCrCard	IsActiveMember
EstimatedSalary \				
0	0.00	1	1	1
101348.88				
1	83807.86	1	0	1
112542.58				
2	159660.80	3	1	0
113931.57				
3	0.00	2	0	0
93826.63				
4	125510.82	1	1	1
79084.10				

	Exited	Geography_France	Geography_Germany	Geography_Spain
0	1	1	0	0
1	0	0	0	1
2	1	1	0	0
3	0	1	0	0
4	0	0	0	1

8. Split the data into dependent and independent variables.

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Gender	Age	
Tenure \							
0	1	15634602	Hargrave	619	0	42	2
1	2	15647311	Hill	608	0	41	1
2	3	15619304	Onio	502	0	42	8
3	4	15701354	Boni	699	0	39	1
4	5	15737888	Mitchell	850	0	43	2

	Balance	NumOfProducts	HasCrCard	IsActiveMember
EstimatedSalary \				
0	0.00	1	1	1
101348.88				
1	83807.86	1	0	1
112542.58				
2	159660.80	3	1	0
113931.57				
3	0.00	2	0	0
93826.63				
4	125510.82	1	1	1
79084.10				

	Exited	Geography_0	Geography_1	Geography_2
0	1	1	0	0
1	0	0	0	1
2	1	1	0	0
3	0	1	0	0
4	0	0	0	1

#Independent variable

```
x=df.iloc[:,[4,5,6,11,12]]
```

```
x.head()
```

	Gender	Age	Tenure	EstimatedSalary	Exited
0	0	42	2	101348.88	1
1	0	41	1	112542.58	0
2	0	42	8	113931.57	1
3	0	39	1	93826.63	0
4	0	43	2	79084.10	0

#Dependent variable

```
y=df.iloc[:,[13]]
```

```
y.head()
```



```

    Geography_0
0             1
1             0
2             1
3             1
4             0

```

9. Scale the independent variables

```

from sklearn.preprocessing import StandardScaler
ss=StandardScaler()

```

```

#x_train = sc.fit_transform(x_train)

```

```

data=ss.fit_transform(x)

```

```

data

```

```

array([[ -1.09598752,  0.29351742, -1.04175968,  0.02188649,
 1.97716468],
       [ -1.09598752,  0.19816383, -1.38753759,  0.21653375, -
0.50577476],
       [ -1.09598752,  0.29351742,  1.03290776,  0.2406869 ,
 1.97716468],
       ...,
       [ -1.09598752, -0.27860412,  0.68712986, -1.00864308,
 1.97716468],
       [  0.91241915,  0.29351742, -0.69598177, -0.12523071,
 1.97716468],
       [ -1.09598752, -1.04143285, -0.35020386, -1.07636976, -
0.50577476]])

```

```

from sklearn.preprocessing import scale
scale=pd.DataFrame(scale(x),columns=x.columns)

```

```

scale

```

	Gender	Age	Tenure	EstimatedSalary	Exited
0	-1.095988	0.293517	-1.041760	0.021886	1.977165
1	-1.095988	0.198164	-1.387538	0.216534	-0.505775
2	-1.095988	0.293517	1.032908	0.240687	1.977165
3	-1.095988	0.007457	-1.387538	-0.108918	-0.505775
4	-1.095988	0.388871	-1.041760	-0.365276	-0.505775
...
9995	0.912419	0.007457	-0.004426	-0.066419	-0.505775
9996	0.912419	-0.373958	1.724464	0.027988	-0.505775
9997	-1.095988	-0.278604	0.687130	-1.008643	1.977165
9998	0.912419	0.293517	-0.695982	-0.125231	1.977165
9999	-1.095988	-1.041433	-0.350204	-1.076370	-0.505775

```

[10000 rows x 5 columns]

```

10. Split the data into training and testing

```
x=df.iloc[:,0:15].values
```

```
y=df.iloc[:,15:16].values
```

```
x.shape
```

```
(10000, 15)
```

```
y.shape
```

```
(10000, 1)
```

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.2,  
random_state=0)
```

```
x_train.shape
```

```
(8000, 15)
```

```
y_train.shape
```

```
(8000, 1)
```